

## 1. Problem and users

Official documents are falsified through edited text, replaced photos, forged stamps, altered QR codes — and now through AI generation. Manual review is slow, repetitive, and misses precisely the subtle cases.

**Who.** Teams that verify documents at volume: recruitment, finance, compliance and administration.

### Positioning

An AI-assisted screening tool that returns **evidence a reviewer can act on** — not a verdict. Every flag is drawn and numbered on the document itself, so the reviewer sees *where* to look, not only *that* something is wrong.

## 2. Evidence

### Tests run against the working prototype

1. AI-generated invoice, visually flawless → scored *low risk*.
2. Client tax identifier of invalid length → missed by the model.
3. Low score returned together with a reject verdict → contradictory output.

Each finding changed the build. Interviews are under way to rank use cases.

## 3. Solution — a division of labour

Upload → Normalization → Multimodal analysis → Deterministic checks → Risk score → Report

### Handled by the model

- ✓ **Reads** the document and extracts fields.
- ✓ **Sees** visual manipulation: mixed fonts, broken alignment, re-inserted text.
- ✓ **Explains** each finding in plain language.

### Never trusted to the model

- ✗ **Counting digits** — identifier length.
- ✗ **Checksums** — IBAN mod-97 key.
- ✗ **Arithmetic** — VAT and column totals.

Handled by deterministic code: exact and auditable.

## 4. Prototype — live

- PDF, JPG, PNG. Preview and model input are the same image, so flagged regions always align.
- Strict JSON schema, re-validated server-side — a malformed analysis is rejected, never half-displayed.
- A failed deterministic check raises the score and blocks an “accept” verdict.

Out of 72h scope, part of the target architecture: OCR tuning, dedicated manipulation models, metadata forensics, cross-document checks.

## 5. Validation and business logic

Focus: how documents are verified today, which frauds are hardest to catch, and what evidence earns a reviewer’s trust. A learning log records who was contacted, what was learned, what was surprising, and what changed.

**Customer** Organizations verifying documents at volume — lenders, insurers, recruiters, schools, administrations.

**Value** Less manual screening, faster triage, traceable evidence.

**Model** B2B SaaS per seat or per document, plus an integration API.

**Market** Finance, recruitment, insurance, education, public sector.

**Next 30 days** Pilot with one organization, add metadata forensics and cross-document consistency, test pricing.

### Key assumption

Organizations will adopt AI-assisted screening only if it shows transparent evidence for every flag. A score without justification will not be trusted — and should not be.