03 Validation Evidence

DocShield AI — tested by a third party, and what changed as a result

1. First external validation — TecForge

Company	TecForge
Tester	Aissam EN-NAHEL, Human Resources — written feedback, 28 August 2026
Material	Three documents: a genuine ISO certificate, a flattened test invoice, and a real two-page invoice the tester altered himself to see whether the tool would notice.
Method	Unassisted use of the deployed prototype, followed by written feedback. No script, no guidance.

Verdict in the tester's own words

The idea was judged useful and the text extraction very good — including on a flattened PDF read as a scanned image. *"Je pense que le score de risque doit être plus fiable"*: the concern was the reliability of the risk score, not the reading of the document.

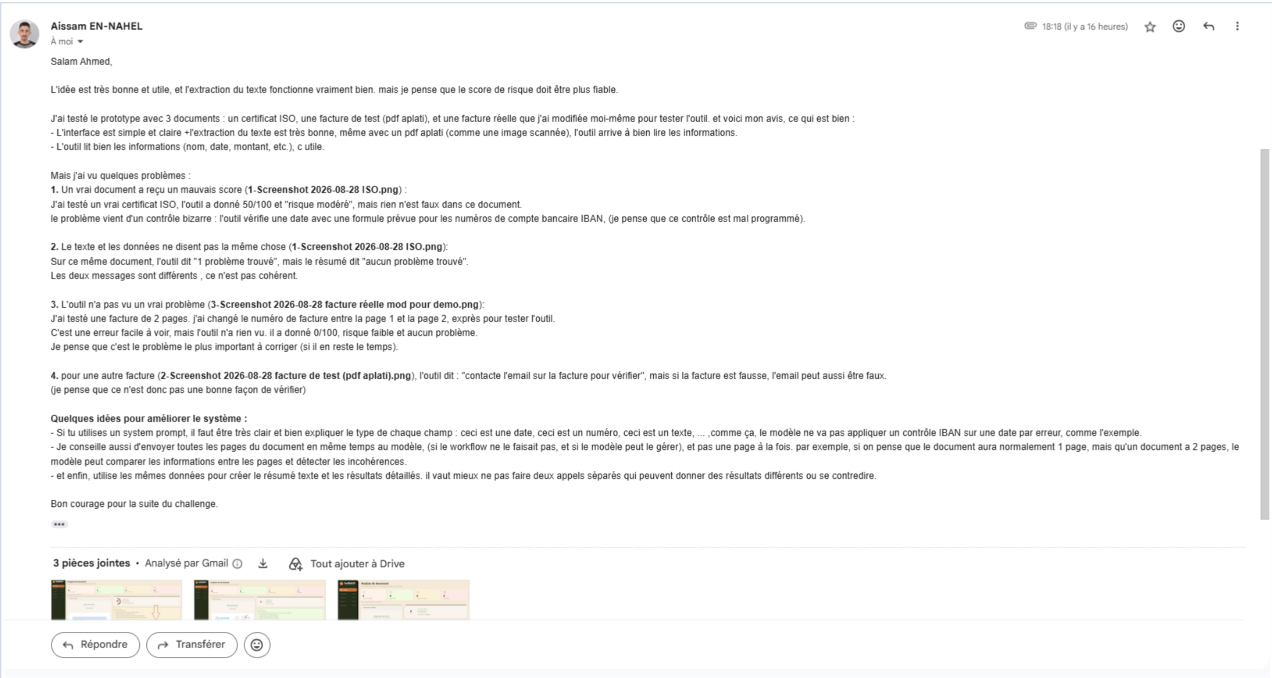
Aissam EN-NAHEL
À moi
18:18 (il y a 16 heures)

Salam Ahmed,

L'idée est très bonne et utile, et l'extraction du texte fonctionne vraiment bien. mais je pense que le score de risque doit être plus fiable.

J'ai testé le prototype avec 3 documents : un certificat ISO, une facture de test (pdf aplati), et une facture réelle que j'ai modifiée moi-même pour tester l'outil. et voici mon avis, ce qui est bien :
- L'interface est simple et claire +l'extraction du texte est très bonne, même avec un pdf aplati (comme une image scannée), l'outil arrive à bien lire les informations.
- L'outil lit bien les informations (nom, date, montant, etc.), c utile.

Mais j'ai vu quelques problèmes :
1. Un vrai document a reçu un mauvais score (**1-Screenshot 2026-08-28 ISO.png**) :
J'ai testé un vrai certificat ISO, l'outil a donné 50/100 et "risque modéré", mais rien n'est faux dans ce document.
le problème vient d'un contrôle bizarre : l'outil vérifie une date avec une formule prévue pour les numéros de compte bancaire IBAN, (je pense que ce contrôle est mal programmé).

2. Le texte et les données ne disent pas la même chose (**1-Screenshot 2026-08-28 ISO.png**):
Sur ce même document, l'outil dit "1 problème trouvé", mais le résumé dit "aucun problème trouvé".
Les deux messages sont différents , ce n'est pas cohérent.

3. L'outil n'a pas vu un vrai problème (**3-Screenshot 2026-08-28 facture réelle mod pour demo.png**):
J'ai testé une facture de 2 pages. j'ai changé le numéro de facture entre la page 1 et la page 2, exprès pour tester l'outil.
C'est une erreur facile à voir, mais l'outil n'a rien vu. il a donné 0/100, risque faible et aucun problème.
Je pense que c'est le problème le plus important à corriger (si il en reste le temps).

4. pour une autre facture (**2-Screenshot 2026-08-28 facture de test (pdf aplati).png**), l'outil dit : "contacte l'email sur la facture pour vérifier", mais si la facture est fausse, l'email peut aussi être faux.
(je pense que ce n'est donc pas une bonne façon de vérifier)

Quelques idées pour améliorer le système :
- Si tu utilises un system prompt, il faut être très clair et bien expliquer le type de chaque champ : ceci est une date, ceci est un numéro, ceci est un texte, ... ,comme ça, le modèle ne va pas appliquer un contrôle IBAN sur une date par erreur, comme l'exemple.
- Je conseille aussi d'envoyer toutes les pages du document en même temps au modèle, (si le workflow ne le faisait pas, et si le modèle peut le gérer), et pas une page à la fois. par exemple, si on pense que le document aura normalement 1 page, mais qu'un document a 2 pages, le modèle peut comparer les informations entre les pages et détecter les incohérences.
- et enfin, utilise les mêmes données pour créer le résumé texte et les résultats détaillés. il vaut mieux ne pas faire deux appels séparés qui peuvent donner des résultats différents ou se contredire.

Bon courage pour la suite du challenge.

3 pièces jointes • Analysé par Gmail
Tout ajouter à Drive
Répondre
Transférer

Feedback received from Aissam EN-NAHEL (TecForge), reproduced verbatim.

Held up under test

- ✓ Interface judged simple and clear.
- ✓ Extraction accurate even on a flattened PDF.
- ✓ Fields read correctly: name, date, amount.

Did not hold up

- × A genuine document scored as suspicious.
- × Narrative contradicted the anomaly count.
- × A deliberate forgery went unnoticed.

2. The four defects, and their cause in the code

Each report was traced to a specific line before anything was changed. None was a tuning problem; all four were logic errors.

1. **A genuine ISO certificate scored 50/100.** The IBAN checksum fired on the shape of a value while ignoring its label. "Two letters, two digits, then 11 to 30 characters" describes a certificate number as well as a bank account, so the mod-97 key failed on an authentic document — and a single failed check imposed a score floor of 50, turning a false positive into "moderate risk".
2. **The count said one problem, the summary said none.** The model writes its explanation before the deterministic checks run. It could therefore report "nothing found" while those checks were adding an anomaly to the list. Two true statements about two different stages, printed as one contradiction.
3. **An invoice number changed between page 1 and page 2 went unseen.** The PDF renderer read `getPage(1)` only, while reporting the real page count. The altered page never reached the model. The tester ranked this the most important defect; the diagnosis confirms it, since no amount of prompt tuning could have recovered a page

that was never sent.

4. **"Contact the email on the invoice to verify."** A forged document carries forged contact details. The prompt asked for "concrete verification steps" without ever forbidding the document from vouching for itself.

3. From feedback to shipped change

Every point raised was corrected and deployed within the challenge window, not deferred to a backlog. The corresponding changes are public in the project repository: github.com/Hajji-ahmed/TrustDoc-AI.

Reported	Change made	Status
Wrong rule applied to a date	The IBAN check now requires a label that names a bank account. An IBAN filed under a label that does not say so escapes the check — accusing a genuine document costs more than missing one verification.	SHIPPED
Summary contradicts detail	Each statement is attributed to its stage: visual analysis first, automatic checks second. The model's prose is kept, not rewritten.	SHIPPED
Two pages never compared	All pages are rendered and sent to the model in one request, each announced by its number. The prompt requires values repeated across pages to be compared.	SHIPPED
Email alone is weak evidence	The prompt forbids any verification routed through the document's own contact details, and requires an independent source instead.	SHIPPED
Field types left implicit *(tester's suggestion)*	The prompt now states which labels may be used for which value types, so the wrong deterministic check cannot be triggered downstream.	SHIPPED
Cross-page detection on the tester's own file	Route contract verified on six cases; the two-page invoice has not been run again through the model.	TO RETEST

What this round actually taught

Extraction quality was never the weak point — the tester confirmed it works, including on flattened scans. Every defect he found sat in the layer that decides *what the reading means*: which rule applies to which field, whether pages are compared, whether the narrative matches the findings. A screening tool is not limited by what it can read, but by the discipline of what it concludes.

The correction that has not yet been proven

Points 1, 2 and 4 are verified by automated assertions. Point 3 — the one the tester ranked highest — is verified only up to the request contract: all pages now reach the model. Whether the model *notices* the altered invoice number requires re-running his own two-page file. Until that is done, this defect is fixed in the pipeline and unproven in the result.

4. Second external validation — Managov

Company	Managov
Reviewer	Abdelkrim Zrouli, Director — written feedback, 28 August 2026
Material	A synthetic test pack he built himself: Document A, a coherent administrative attestation, and Document B, the same document with its issue date made inconsistent with its validity period.
Method	Unprompted assessment of the concept, followed by a controlled A/B test on a matched pair of documents.

Assessment

“DocShield AI répond à un besoin réel de pré-contrôle documentaire.” The combination of AI analysis, confidence score, anomaly detection and recommendation was judged relevant, the interface simple and intuitive, and the project promising for public administrations, educational institutions and companies.

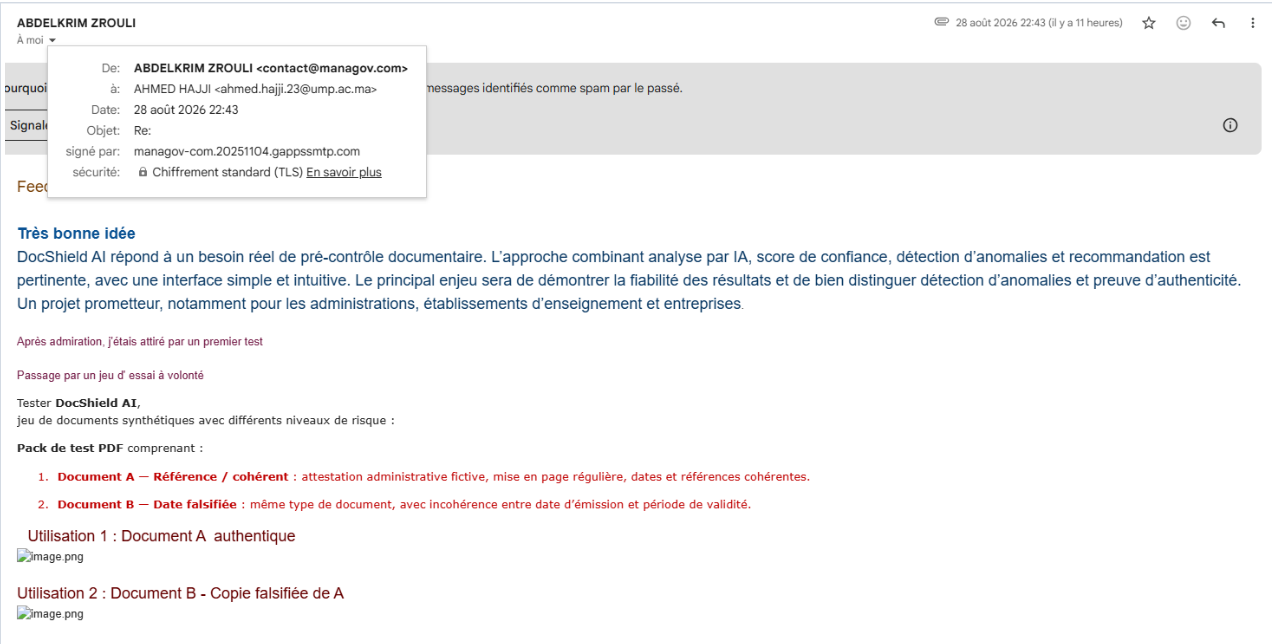

ABDELKRIM ZROULI
À moi

28 août 2026 22:43 (il y a 11 heures)

De: ABDELKRIM ZROULI <contact@managov.com>
à: AHMED HAJJI <ahmed.hajji.23@ump.ac.ma>
Date: 28 août 2026 22:43
Objet: Re:
signé par: managov-com.20251104.gappssmtp.com
sécurité: Chiffrement standard (TLS) En savoir plus

ourquoi … messages identifiés comme spam par le passé.

Signal

Feed

Très bonne idée

DocShield AI répond à un besoin réel de pré-contrôle documentaire. L’approche combinant analyse par IA, score de confiance, détection d’anomalies et recommandation est pertinente, avec une interface simple et intuitive. Le principal enjeu sera de démontrer la fiabilité des résultats et de bien distinguer détection d’anomalies et preuve d’authenticité. Un projet prometteur, notamment pour les administrations, établissements d’enseignement et entreprises.

Après admiration, j'étais attiré par un premier test

Passage par un jeu d' essai à volonté

Tester **DocShield AI**,
jeu de documents synthétiques avec différents niveaux de risque :

Pack de test PDF comprenant :

1. **Document A — Référence / cohérent** : attestation administrative fictive, mise en page régulière, dates et références cohérentes.
2. **Document B — Date falsifiée** : même type de document, avec incohérence entre date d’émission et période de validité.

Utilisation 1 : Document A authentique
image.png

Utilisation 2 : Document B - Copie falsifiée de A
image.png

Feedback received from Abdelkrim Zrouli (Managov), reproduced verbatim.

The reservation, and why it matters more than the praise

“Le principal enjeu sera de démontrer la fiabilité des résultats et de bien distinguer détection d’anomalies et preuve d’authenticité.”

This is the product’s own founding constraint, stated back by someone who had not been told it. DocShield never claims a document is authentic; it reports what looks irregular and hands the decision to a human. Hearing that boundary named as *the* central challenge, by a director evaluating the tool for his own organisation, confirms that the distinction is not academic — it is what buyers will judge the product on.

4.1 The test design — one controlled variable

Both documents are byte-for-byte identical except for the validity window. That is a cleaner experiment than anything run in-house: it isolates a single failure mode and ships with its own control.

Field	Document A — control	Document B — altered
Issue date	28 August 2026	28 August 2026
Period covered	1 Jan 2026 to 31 Dec 2026	1 Jan 2026 to **31 Dec 2025**
Validity	until 31 Dec 2026	until **31 Dec 2025**
Everything else	Identical — reference, identifier, issuer, signature, verification reference	

The planted defect is a period that ends before it begins, on a document issued after both dates. It is obvious to a human and invisible to any check that reads fields one at a time.

4.2 The result — the planted defect was caught

Document B, scored by the deployed prototype

Score 34/100 — moderate risk. **Type** Attestation administrative, correctly identified.

Anomaly One, and only one: *"La période concernée commence en 2026 mais termine en 2025, ce qui est chronologiquement impossible."* Classified under data coherence, moderate severity, and drawn as a numbered region on the document itself.

Recommendation Manual verification required — confirm the details with the issuer *through alternative official channels*, and compare against other documents.

What this result confirms beyond the detection itself

Two of the corrections made after the first validation are visible in this output, independently produced by a second reviewer who knew nothing of them.

The narrative no longer contradicts the findings. The explanation reads *"Analyse visuelle: . . . "* then *"Contrôles automatiques: aucun écart relevé. . . "*. Each statement is attributed to its stage, and both are true at once: the model saw the date problem, the deterministic checks correctly found nothing in their own scope.

The verification step no longer routes through the document. The recommendation names *alternative official channels* rather than the contact details printed on the page — the precise flaw reported by the first reviewer.

The control was not scored

Document A exists to answer the question the reviewer wrote at the top of both files: *"Vérifier que DocShield AI ne génère pas de faux positifs."* It was uploaded, but the captured screen shows it before the analysis was run, so **no result is recorded for the control**.

Detection on B is therefore established; the absence of a false positive on A is not. A test that only ever fires proves nothing on its own — the control is what separates a detector from an alarm that is always on. Running Document A is the single most useful next step in this dossier.

5. Third external validation — INKWAY

Company	INKWAY
Reviewer	Chaimae EL OUAZZANI, Project Manager — written feedback, 29 August 2026
Material	No document test. An assessment of the product direction, given once the current scope had been explained.
Method	Unprompted review, returned as eight numbered proposals, each with the operational situation it would serve.

Where the idea earns its place

The concept was judged useful *"surtout si elle est appliquée à des situations où les organisations doivent contrôler beaucoup de documents en peu de temps"*. Volume under time pressure is the qualifying condition — the second reviewer to reach that criterion independently, and the same reasoning that put lending and leasing at the head of deliverable 04.

5.1 The eight proposals, against what exists

Recorded as a roadmap rather than a wish list, each measured against the prototype as it stands today.

Proposed	Against the prototype	Status
Diplomas and certificates: QR code, certificate number, issuing body	Internal coherence is checked. Reading a QR code and looking up an issuer are not.	NOT BUILT
Reference attestations: dates, amounts, companies, existence of the issuer	Dates and arithmetic are checked. Confirming that the issuer exists requires an external source.	PARTIAL
Comparison against known document templates	The model judges plausibility, never conformity to a reference model. It has no template to compare with.	NOT BUILT
Change of bank details or other sensitive fields	The IBAN check key is verified. Detecting a *change* requires knowing the supplier's usual details — a history the tool does not keep.	PARTIAL
Signatures and stamps: copied, pasted, added after the fact	The model is asked to inspect seals and signatures visually. No forensic analysis.	PARTIAL
Metadata and file history	Excluded by construction: the document is rasterised before analysis, so metadata never reaches the model.	NOT BUILT
Cross-checking several documents of one file	Cross-*page* comparison within a single document ships today. Cross-*document* does not.	PARTIAL
Connection to external sources	Not built. This is the boundary the second reviewer named, reached here from a different direction.	NOT BUILT

The reframing, which matters more than any single item

"Ne pas seulement analyser si le document a l'air correct, mais aider une organisation à sécuriser une décision réelle" — awarding a contract, granting credit, hiring, paying a supplier, settling a claim.

That moves the unit of value from the document to the decision it supports. Six of the eight proposals follow from it: they are not detection features, they are what a reviewer needs before signing. It also explains why external sources recur — a decision cannot be secured by a document vouching for itself.

On the interface The reviewer would replace any wording suggesting *authentic* with *low risk / to be checked / high risk*, plus the precise reason for the flag. The risk badge already reads that way; the recommendation label *"Document acceptable"* does not, and still implies a verdict on authenticity. Small change, and directly on the boundary all three reviewers have now named.

What three independent reviews converge on

An HR practitioner, a company director and a project manager, none aware of the others' feedback, arrived at the same boundary from three directions: the score must be trustworthy, detection must stay distinct from proof, and the interface must never say *authentic*.

No reviewer asked for more detection categories. All three asked for the product to be honest about what it does not establish. That is the clearest product signal the challenge produced.

6. Validation register

One line per company contacted during the challenge. Only verified conversations and real feedback are recorded here; the blank rows are reserved for the validations still in progress.

Company	Person / Role	Feedback / Learning	Change / Next action
TecForge	Aissam EN-NAHEL *Human Resources*	Document reading judged reliable, including on flattened scans; the risk assessment was not. Four defects reported: a validation rule applied to the wrong type of field, a summary inconsistent with the detailed findings, no comparison between the pages of a single document, and a verification step relying on the document's own contact details.	All four defects corrected and deployed within the challenge window. Next: retest the altered two-page invoice, then extend testing to other document families.
Managov	Abdelkrim Zrouli *Director*	Concept and interface judged relevant for administrations, schools and companies. Central challenge named: demonstrating the reliability of results, and keeping anomaly detection clearly distinct from proof of authenticity. Ran a matched pair of synthetic documents differing by one falsified validity window; the altered document was correctly flagged.	Positioning confirmed and kept explicit in the interface. Detection validated on the altered document. Next: run the control document to confirm no false positive is raised.
INKWAY	Chaimae EL OUAZZANI *Project Manager*	Concept judged useful where organisations must check many documents under time pressure. Eight proposals returned, each tied to an operational situation, all pointing the same way: secure a decision, not score a document. Would drop any wording suggesting authenticity in favour of a risk level with its reason.	Proposals recorded against the current state as a roadmap. Interface wording to revise. Volume under time pressure confirmed as the qualifying condition.

7. Summary and evidence checklist

Completed Three external companies, three reviewers — an HR practitioner, a company director and a project manager — seven documents tested, four defects found and four corrections shipped within the challenge window.

Confirmed from outside The boundary the product was built on — report anomalies, never certify authenticity — was reached independently by all three reviewers, none of them told it and none aware of the others. Two of the corrections made after the first validation are visible in the second reviewer's test output. No reviewer asked for more detection categories.

Detection demonstrated On a falsification designed by a third party, the prototype isolated the planted defect, named it correctly and localised it on the document.

Not yet demonstrated The absence of false positives. The control document of that same test was never scored, and the two-page invoice from the first validation has not been re-run.

Next step Run Document A, the control. It is the cheapest test remaining and the only one that turns the current evidence into a complete result.

Evidence checklist

✓ Three companies contacted, all three replied with written feedback
✓ Root cause identified in code for all four defects
✓ Corrections implemented and deployed, published in the project repository
✓ Detection confirmed on a third party's falsified document
✓ Eight product proposals recorded against the current state
× Third message not yet captured as a screenshot
× Control document not scored — absence of false positives unverified
× Retest of the two-page invoice with the corrected prototype