03 Validation Evidence

DocShield AI — tested by a third party, and what changed as a result

1. First external validation — TecForge

Company	TecForge
Tester	Aissam EN-NAHEL, Human Resources — written feedback, 28 August 2026
Material	Three documents: a genuine ISO certificate, a flattened test invoice, and a real two-page invoice the tester altered himself to see whether the tool would notice.
Method	Unassisted use of the deployed prototype, followed by written feedback. No script, no guidance.

Verdict in the tester's own words

The idea was judged useful and the text extraction very good — including on a flattened PDF read as a scanned image. *"Je pense que le score de risque doit être plus fiable"*: the concern was the reliability of the risk score, not the reading of the document.

L'idée est très bonne et utile, et l'extraction du texte fonctionne vraiment bien. mais je pense que le score de risque doit être plus fiable.

J'ai testé le prototype avec 3 documents : un certificat ISO, une facture de test (pdf aplati), et une facture réelle que j'ai modifiée moi-même pour tester l'outil. et voici mon avis, ce qui est bien :
- L'interface est simple et claire +l'extraction du texte est très bonne, même avec un pdf aplati (comme une image scannée), l'outil arrive à bien lire les informations.
- L'outil lit bien les informations (nom, date, montant, etc.), c utile.

Mais j'ai vu quelques problèmes :
1. Un vrai document a reçu un mauvais score (**1-Screenshot 2026-08-28 ISO.png**) :
J'ai testé un vrai certificat ISO, l'outil a donné 50/100 et "risque modéré", mais rien n'est faux dans ce document.
le problème vient d'un contrôle bizarre : l'outil vérifie une date avec une formule prévue pour les numéros de compte bancaire IBAN, (je pense que ce contrôle est mal programmé).

2. Le texte et les données ne disent pas la même chose (**1-Screenshot 2026-08-28 ISO.png**):
Sur ce même document, l'outil dit "1 problème trouvé", mais le résumé dit "aucun problème trouvé".
Les deux messages sont différents , ce n'est pas cohérent.

3. L'outil n'a pas vu un vrai problème (**3-Screenshot 2026-08-28 facture réelle mod pour demo.png**):
J'ai testé une facture de 2 pages. j'ai changé le numéro de facture entre la page 1 et la page 2, exprès pour tester l'outil.
C'est une erreur facile à voir, mais l'outil n'a rien vu. il a donné 0/100, risque faible et aucun problème.
Je pense que c'est le problème le plus important à corriger (si il en reste le temps).

4. pour une autre facture (**2-Screenshot 2026-08-28 facture de test (pdf aplati).png**), l'outil dit : "contacte l'email sur la facture pour vérifier", mais si la facture est fausse, l'email peut aussi être faux.
(je pense que ce n'est donc pas une bonne façon de vérifier)

Quelques idées pour améliorer le système :
- Si tu utilises un system prompt, il faut être très clair et bien expliquer le type de chaque champ : ceci est une date, ceci est un numéro, ceci est un texte, ... ,comme ça, le modèle ne va pas appliquer un contrôle IBAN sur une date par erreur, comme l'exemple.
- Je conseille aussi d'envoyer toutes les pages du document en même temps au modèle, (si le workflow ne le faisait pas, et si le modèle peut le gérer), et pas une page à la fois. par exemple, si on pense que le document aura normalement 1 page, mais qu'un document a 2 pages, le modèle peut comparer les informations entre les pages et détecter les incohérences.
- et enfin, utilise les mêmes données pour créer le résumé texte et les résultats détaillés. il vaut mieux ne pas faire deux appels séparés qui peuvent donner des résultats différents ou se contredire.

Bon courage pour la suite du challenge.

Feedback received from Aissam EN-NAHEL (TecForge), reproduced verbatim.

Held up under test

- ✓ Interface judged simple and clear.
- ✓ Extraction accurate even on a flattened PDF.
- ✓ Fields read correctly: name, date, amount.

Did not hold up

- × A genuine document scored as suspicious.
- × Narrative contradicted the anomaly count.
- × A deliberate forgery went unnoticed.

2. The four defects, and their cause in the code

Each report was traced to a specific line before anything was changed. None was a tuning problem; all four were logic errors.

1. **A genuine ISO certificate scored 50/100.** The IBAN checksum fired on the shape of a value while ignoring its label. "Two letters, two digits, then 11 to 30 characters" describes a certificate number as well as a bank account, so the mod-97 key failed on an authentic document — and a single failed check imposed a score floor of 50, turning a false positive into "moderate risk".
2. **The count said one problem, the summary said none.** The model writes its explanation before the deterministic checks run. It could therefore report "nothing found" while those checks were adding an anomaly to the list. Two true statements about two different stages, printed as one contradiction.
3. **An invoice number changed between page 1 and page 2 went unseen.** The PDF renderer read `getPage(1)` only, while reporting the real page count. The altered page never reached the model. The tester ranked this the most important defect; the diagnosis confirms it, since no amount of prompt tuning could have recovered a page that was never sent.
4. **"Contact the email on the invoice to verify."** A forged document carries forged contact details. The prompt asked for "concrete verification steps" without ever forbidding the document from vouching for itself.

3. From feedback to shipped change

Every point raised was fixed and pushed, not queued. Commit `761c596` on github.com/Hajji-ahmed/TrustDoc-AI.

Reported	Change made	Status
Wrong rule applied to a date	The IBAN check now requires a label that names a bank account. An IBAN filed under a label that does not say so escapes the check — accusing a genuine document costs more than missing one verification.	SHIPPED
Summary contradicts detail	Each statement is attributed to its stage: visual analysis first, automatic checks second. The model's prose is kept, not rewritten.	SHIPPED
Two pages never compared	All pages are rendered and sent to the model in one request, each announced by its number. The prompt requires values repeated across pages to be compared.	SHIPPED
Email alone is weak evidence	The prompt forbids any verification routed through the document's own contact details, and requires an independent source instead.	SHIPPED
Field types left implicit *(tester's suggestion)*	The prompt now states which labels may be used for which value types, so the wrong deterministic check cannot be triggered downstream.	SHIPPED
Cross-page detection on the tester's own file	Route contract verified on six cases; the two-page invoice has not been run again through the model.	TO RETEST

What this round actually taught

Extraction quality was never the weak point — the tester confirmed it works, including on flattened scans. Every defect he found sat in the layer that decides *what the reading means*: which rule applies to which field, whether pages are compared, whether the narrative matches the findings. A screening tool is not limited by what it can read, but by the discipline of what it concludes.

The correction that has not yet been proven

Points 1, 2 and 4 are verified by automated assertions. Point 3 — the one the tester ranked highest — is verified only up to the request contract: all pages now reach the model. Whether the model *notices* the altered invoice number requires re-running his own two-page file. Until that is done, this defect is fixed in the pipeline and unproven in the result.

4. Validation register

One line per company contacted during the challenge. Only verified conversations and real feedback are recorded here; the blank rows are reserved for the validations still in progress.

Company	Person / Role	Feedback / Learning	Change / Next action
TecForge	Aissam EN-NAHEL *Human Resources*	Extraction is reliable; the risk score is not. Four defects: a rule applied to the wrong field type, the summary contradicting the detail, pages never compared, and verification routed through the document itself.	Four corrections shipped, commit `761c596`. Retest the altered two-page invoice, then extend to other document families.

5. Summary and evidence checklist

Completed One external company, one detailed tester, three documents, four defects found and four corrections shipped within the challenge window.

Highest-priority correction Detecting inconsistencies between pages of the same document — shipped, awaiting retest on the tester's file.

Next step Re-run the altered two-page invoice, then widen testing to the companies listed above.

Evidence checklist

- ✓ TecForge feedback captured, reproduced verbatim
- ✓ Screenshot evidence included
- ✓ Root cause identified in code for all four defects
- ✓ Corrections implemented and pushed — commit `761c596`
- × Retest of the two-page invoice with the corrected prototype
- × Additional companies, feedback and resulting changes