


1. First external validation — TecForge

Company	TecForge
Tester	Aissam EN-NAHEL, Human Resources — written feedback, 28 August 2026
Material	Three documents: a genuine ISO certificate, a flattened test invoice, and a real two-page invoice the tester altered himself to see whether the tool would notice.
Method	Unassisted use of the deployed prototype, followed by written feedback. No script, no guidance.

Verdict in the tester's own words

The idea was judged useful and the text extraction very good — including on a flattened PDF read as a scanned image. *“Je pense que le score de risque doit être plus fiable”*: the concern was the reliability of the risk score, not the reading of the document.



Feedback received from Aissam EN-NAHEL (TecForge), reproduced verbatim.

Held up under test

- ✓ Interface judged simple and clear.
- ✓ Extraction accurate even on a flattened PDF.
- ✓ Fields read correctly: name, date, amount.

Did not hold up

- ✗ A genuine document scored as suspicious.
- ✗ Narrative contradicted the anomaly count.
- ✗ A deliberate forgery went unnoticed.

2. The four defects, and their cause in the code

Each report was traced to a specific line before anything was changed. None was a tuning problem; all four were logic errors.

- 1. A genuine ISO certificate scored 50/100.** The IBAN checksum fired on the shape of a value while ignoring its label. “Two letters, two digits, then 11 to 30 characters” describes a certificate number as well as a bank account, so the mod-97 key failed on an authentic document — and a single failed check imposed a score floor of 50, turning a false positive into “moderate risk”.
- 2. The count said one problem, the summary said none.** The model writes its explanation before the deterministic checks run. It could therefore report “nothing found” while those checks were adding an anomaly to the list. Two true statements about two different stages, printed as one contradiction.
- 3. An invoice number changed between page 1 and page 2 went unseen.** The PDF renderer read `getPage(1)` only, while reporting the real page count. The altered page never reached the model. The tester ranked this the most important defect; the diagnosis confirms it, since no amount of prompt tuning could have recovered a page

that was never sent.

4. **“Contact the email on the invoice to verify.”** A forged document carries forged contact details. The prompt asked for “concrete verification steps” without ever forbidding the document from vouching for itself.

3. From feedback to shipped change

Every point raised was corrected and deployed within the challenge window, not deferred to a backlog. The corresponding changes are public in the project repository: github.com/Hajji-ahmed/TrustDoc-AI.

Reported	Change made	Status
Wrong rule applied to a date	The IBAN check now requires a label that names a bank account. An IBAN filed under a label that does not say so escapes the check — accusing a genuine document costs more than missing one verification.	SHIPPED
Summary contradicts detail	Each statement is attributed to its stage: visual analysis first, automatic checks second. The model’s prose is kept, not rewritten.	SHIPPED
Two pages never compared	All pages are rendered and sent to the model in one request, each announced by its number. The prompt requires values repeated across pages to be compared.	SHIPPED
Email alone is weak evidence	The prompt forbids any verification routed through the document’s own contact details, and requires an independent source instead.	SHIPPED
Field types left implicit (tester’s suggestion)	The prompt now states which labels may be used for which value types, so the wrong deterministic check cannot be triggered downstream.	SHIPPED
Cross-page detection on the tester’s own file	Route contract verified on six cases; the two-page invoice has not been run again through the model.	TO RETEST

What this round actually taught

Extraction quality was never the weak point — the tester confirmed it works, including on flattened scans. Every defect he found sat in the layer that decides *what the reading means*: which rule applies to which field, whether pages are compared, whether the narrative matches the findings. A screening tool is not limited by what it can read, but by the discipline of what it concludes.

The correction that has not yet been proven


Points 1, 2 and 4 are verified by automated assertions. Point 3 — the one the tester ranked highest — is verified only up to the request contract: all pages now reach the model. Whether the model *notices* the altered invoice number requires re-running his own two-page file. Until that is done, this defect is fixed in the pipeline and unproven in the result.

4. Second external validation — Managov

Company	Managov
Reviewer	Abdelkrim Zrouli, Director — written feedback, 28 August 2026
Material	A synthetic test pack he built himself: Document A, a coherent administrative attestation, and Document B, the same document with its issue date made inconsistent with its validity period.
Method	Unprompted assessment of the concept, followed by a controlled A/B test on a matched pair of documents.

Assessment

“DocShield AI répond à un besoin réel de pré-contrôle documentaire.” The combination of AI analysis, confidence score, anomaly detection and recommendation was judged relevant, the interface simple and intuitive, and the project promising for public administrations, educational institutions and companies.



Feedback received from Abdelkrim Zrouli (Managov), reproduced verbatim.

The reservation, and why it matters more than the praise

“Le principal enjeu sera de démontrer la fiabilité des résultats et de bien distinguer détection d’anomalies et preuve d’authenticité.”

This is the product’s own founding constraint, stated back by someone who had not been told it. DocShield never claims a document is authentic; it reports what looks irregular and hands the decision to a human. Hearing that boundary named as *the* central challenge, by a director evaluating the tool for his own organisation, confirms that the distinction is not academic — it is what buyers will judge the product on.

Test designed better than ours — and not yet scored

Document A and Document B differ by one controlled variable: an inconsistency between issue date and validity period. That is a cleaner experiment than anything run in-house, because it isolates a single failure mode and provides its own control.

The two result screenshots did not render in the received message, so the outcome of this A/B test is **not documented here**. Reporting it from memory or inference would defeat the purpose of an evidence document. The result is being requested from the reviewer, and the pair will be re-run against the corrected prototype in either case.

5. Validation register

One line per company contacted during the challenge. Only verified conversations and real feedback are recorded here; the blank rows are reserved for the validations still in progress.

Company	Person / Role	Feedback / Learning	Change / Next action
TecForge	Aissam EN-NAHEL <i>Human Resources</i>	Document reading judged reliable, including on flattened scans; the risk assessment was not. Four defects reported: a validation rule applied to the wrong type of field, a summary inconsistent with the detailed findings, no comparison between the pages of a single document, and a verification step relying on the document's own contact details.	All four defects corrected and deployed within the challenge window. Next: retest the altered two-page invoice, then extend testing to other document families.
Managov	Abdelkrim Zrouli <i>Director</i>	Concept and interface judged relevant for administrations, schools and companies. Central challenge named: demonstrating the reliability of results, and keeping anomaly detection clearly distinct from proof of authenticity. Ran a matched pair of synthetic documents differing by one falsified date.	Positioning confirmed and kept explicit in the interface. Next: obtain the two result screenshots, then re-run the pair against the corrected prototype.

6. Summary and evidence checklist

Completed Two external companies, two reviewers — one HR practitioner and one director — five documents tested, four defects found and four corrections shipped within the challenge window.

Confirmed from outside The boundary the product was built on — report anomalies, never certify authenticity — was independently named as the central challenge by the second reviewer, who had not been told it.

Highest-priority correction Detecting inconsistencies between pages of the same document — shipped, awaiting retest on the reviewers' own files.

Next step Obtain the Managov result screenshots, re-run both the altered two-page invoice and the A/B pair, then widen testing further.

Evidence checklist

- ✓ Two companies contacted, both replied with written feedback
- ✓ Both messages captured and reproduced verbatim
- ✓ Root cause identified in code for all four defects
- ✓ Corrections implemented and deployed, published in the project repository
- ✗ Managov A/B test results — screenshots did not render in the message
- ✗ Retest of the two-page invoice and of the A/B pair with the corrected prototype