

04 Business Logic

DocShield AI — customer, value, model, market, next step

1. Customer — who pays, and who signs

User The document reviewer: a credit analyst, an HR recruiter, a compliance officer. They feel the pain but rarely hold the budget.

Buyer The head of risk, compliance or operations. They do not buy time savings; they buy *reduced exposure* — a forged payslip approved, a diploma never checked, a regulator asking how a file was verified.

Beachhead Lending and leasing companies in Morocco. High document volume, direct financial loss on a miss, and a verification step that already exists — so the product replaces a manual task rather than creating a new one.

2. What two external reviews changed

The prototype was put in front of two people during the challenge, chosen for opposite vantage points: an HR practitioner who verifies documents himself, and a company director who would sign for the tool. Deliverable 03 records the detail; what follows is only what bears on the business.

The user tested the mechanism. Reading and extraction were judged reliable, including on a flattened scan. The objection was not that the tool fails to see — it was that the score could not yet be trusted. For a screening product, that is the whole sale.

The buyer named the criterion. Unprompted, the director identified the central challenge as *demonstrating the reliability of the results, and keeping anomaly detection clearly distinct from proof of authenticity*. He was describing the product's founding constraint without having been told it.

The beachhead was not confirmed — it was redirected

Section 1 assumes lending and leasing. The one director asked pointed instead to **public administrations, educational institutions and companies**, on his own initiative.

One opinion does not move a beachhead: a single data point cannot outweigh the reasoning behind lending, which rests on direct financial loss per miss. But it is the only segment signal the challenge produced from someone evaluating the tool for his own organisation, and it is recorded here rather than discarded for disagreeing with the hypothesis. It is the first question for the ten interviews in week 1.

What this reorders. Both reviewers converged on the same purchase criterion, and it is not detection coverage. It is whether a flag can be trusted, and whether the product is honest about what it does not prove. That puts the audit trail and the calibration evidence ahead of new detection categories on the roadmap — a ranking neither reviewer was asked about, and both implied.

3. Value — what is actually being bought

Measurable

Verification time per file. Share of files needing a second reviewer. A written, timestamped trace of every check performed.

Not yet measured

Minutes saved per document, and the annual cost of an undetected forgery. Both must come from interviews, not from an assumption.

The second column is deliberately empty of numbers. Quantifying the value before speaking to a single reviewer would produce a convincing figure and a worthless one.

4. Business model

Tier	For	Basis
Pay per document	Low volume, occasional checks	Per analysis, prepaid credits
Per seat	Teams with a recurring flow	Monthly, per reviewer, fair-use cap
API & integration	Case-management systems	Annual contract, volume tiers, SLA

Unit economics — the one number already known

Each analysis costs a few cents of model inference, readable directly from the provider's usage dashboard. Cost scales per document, not per customer, and falls as models get cheaper. Gross margin is therefore high and predictable; the real cost centres are sales and the trust-building work, not compute.

5. Market — sized bottom-up, not asserted

Rather than quote a market figure, the opportunity is expressed as a formula whose inputs are verifiable one by one:

Organizations in segment × reviewers per organization × price per seat × 12

[hypothesis] Each input is a research task, not a guess: counts come from professional registers, reviewer headcount from interviews, price from the first pilot negotiation. *The formula is filled in only once those three are known.*

Expansion path Lending → insurance → recruitment → education → public administration. Same engine, new document families and new business rules — which is also the natural order of increasing sales-cycle length.

6. Next 30 days

Week 1	Ten interviews with reviewers and their managers. Measure the current process before proposing to change it.
Week 2	Close the production gaps: rate limiting, payload ceiling, timestamped audit report. Without them there is no pilot.
Week 3	Run one organization's real files through the tool, side by side with their manual check. Compare, and record every disagreement.
Week 4	Put a price on the table and see whether anyone pays. Fill in the market formula with the three inputs now known.

What would break this

If reviewers do not trust a flag they cannot verify themselves, the tool adds a step instead of removing one. That is the failure mode to test first — before detection accuracy, before pricing, before market size.