04 Business Logic

DocShield AI — customer, value, model, market, next step

1. Customer — who pays, and who signs

User The document reviewer: a credit analyst, an HR recruiter, a compliance officer. They feel the pain but rarely hold the budget.

Buyer The head of risk, compliance or operations. They do not buy time savings; they buy *reduced exposure* — a forged payslip approved, a diploma never checked, a regulator asking how a file was verified.

Beachhead (hypothesis — not yet confirmed) Lending and leasing companies in Morocco: high document volume, direct financial loss on a miss, and a verification step that already exists, so the product replaces a manual task rather than creating one. Two reviewers pointed instead to public administration and education. The hypothesis stands because of the loss-per-miss argument, not because anyone has confirmed it.

2. What two external reviews changed

The prototype was put in front of two people during the challenge, chosen for opposite vantage points: an HR practitioner who verifies documents himself, and a company director who would sign for the tool. Deliverable 03 records the detail; what follows is only what bears on the business.

The user tested the mechanism. Reading and extraction were judged reliable, including on a flattened scan. The objection was not that the tool fails to see — it was that the score could not yet be trusted. For a screening product, that is the whole sale.

The buyer named the criterion. Unprompted, the director identified the central challenge as *demonstrating the reliability of the results, and keeping anomaly detection clearly distinct from proof of authenticity*. He was describing the product's founding constraint without having been told it.

The beachhead was not confirmed — it was redirected

Section 1 assumes lending and leasing. The one director asked pointed instead to **public administrations, educational institutions and companies**, on his own initiative.

One opinion does not move a beachhead: a single data point cannot outweigh the reasoning behind lending, which rests on direct financial loss per miss. But it is the only segment signal the challenge produced from someone evaluating the tool for his own organisation, and it is recorded here rather than discarded for disagreeing with the hypothesis. It is the first question for the ten interviews in week 1.

What this reorders. Both reviewers converged on the same purchase criterion, and it is not detection coverage. It is whether a flag can be trusted, and whether the product is honest about what it does not prove. That puts the audit trail and the calibration evidence ahead of new detection categories on the roadmap — a ranking neither reviewer was asked about, and both implied.

3. Value — what is actually being bought

Measurable	Not yet measured
Verification time per file. Share of files needing a second reviewer. A written, timestamped trace of every check performed.	Minutes saved per document, and the annual cost of an undetected forgery. Both must come from interviews, not from an assumption.

The second column is deliberately empty of numbers. Quantifying the value before speaking to a single reviewer would produce a convincing figure and a worthless one.

What will be measured, in this order Verification time per document; share of files sent to a second reviewer; false positive rate; false negative rate; completeness of the audit trail. The first two size the value. The next two decide whether anyone can rely on it — and both external tests so far have touched only one side of that pair.

4. Business model (hypothesis)

Three tiers, none of them tested. Prices and packaging are to be validated in pilot and pricing conversations; nothing below has been offered to a customer.

Tier	For	Basis
Pay per document	Low volume, occasional checks	Per analysis, prepaid credits
Per seat	Teams with a recurring flow	Monthly, per reviewer, fair-use cap
API & integration	Case-management systems	Annual contract, volume tiers, SLA

Unit economics — the one number already known

Each analysis costs a few cents of model inference, readable directly from the provider's usage dashboard. Cost scales per document, not per customer, and falls as models get cheaper. Gross margin is therefore high and predictable; the real cost centres are sales and the trust-building work, not compute.

5. Market — sized bottom-up, not asserted

Rather than quote a market figure, the opportunity is expressed as a formula whose inputs are verifiable one by one:

Organizations in segment × reviewers per organization × price per seat × 12

[hypothesis] Each input is a research task, not a guess: counts come from professional registers, reviewer headcount from interviews, price from the first pilot negotiation. *The formula is filled in only once those three are known.*

Expansion path Lending → insurance → recruitment → education → public administration. Same engine, new document families and new business rules — which is also the natural order of increasing sales-cycle length.

6. Next 30 days

Week 1	Ten interviews with reviewers and their managers, and a controlled test set: matched genuine and altered pairs, built the way the second reviewer built his. Measure the current process before proposing to change it.
Week 2	Reliability. Run the outstanding retests, then measure false positives and false negatives on that set. Make every flag explainable in one line. Nothing new is built until these numbers exist.
Week 3	Pilot. One organization's real files through the tool, side by side with their manual check. Compare, and record every disagreement — the disagreements are the finding.
Week 4	Price. Put a number on the table and see whether anyone pays. Fill in the market formula with the three inputs then known.

The order is deliberate and follows the risks the validation exposed: trust first, error rates second, real files third, price last. Not one new detection category appears before reliability is measured — a reviewer proposed eight, and they wait.

What would break this

If reviewers do not trust a flag they cannot verify themselves, the tool adds a step instead of removing one. That is the failure mode to test first — before detection accuracy, before pricing, before market size.